

GridSense AI

Explainable Time-Series Intelligence & Anomaly Detection Layer for BESCOM

Hackathon	AI for Bharat
Theme	Theme 8: AI for Smart Meter Intelligence & Loss Detection by BESCOM
Project Title	GridSense AI
Phase	Prototype Phase (Round 2)
Team	Team 007
Live Demo	https://huggingface.co/spaces/rishav123987/gridsense-ai-smart-meter-dashboard
Source Code	https://github.com/lokesh07kewat/BESCOM-Smart-Intelligence
Tech Stack	Python 3.13 scikit-learn Streamlit Plotly SHAP Docker(on huggingface)

All data in this project is synthetic. No real BESCOM customer data is used or stored.

1. Executive Summary

BESCOM's smart meter infrastructure generates over 432,000 readings every 90 days across its urban feeder network, yet this data largely sits unused beyond billing. **GridSense AI** transforms this passive data layer into an active, always-on intelligence system — operating as a read-only decision-support overlay that requires **zero modifications to existing BESCOM infrastructure**.

The system ingests 15-minute interval consumption data and runs two parallel AI pipelines: **Part A** forecasts short-term electricity demand at zone/feeder level and classifies each zone into risk tiers before peak load events occur; **Part B** detects electricity theft, meter tampering, sudden drops, and peer-group deviations — with full explainability for every flag raised.

Every output is accompanied by a plain-English explanation, a SHAP feature breakdown, and a full audit trail. **Nothing is a black box.** Every alert that reaches an inspector is one that a human can understand, verify, and act on independently.

2. Problem Statement

2.1 Demand Side — Reactive Grid Operations

Feeder-level demand spikes are currently unpredictable. Load shedding decisions are made reactively after stress has already occurred. Without short-term localised forecasts, grid operations cannot pre-position resources, alert field teams, or manage transformer loads before overload. This results in avoidable outages and equipment wear.

2.2 Loss Side — Manual Theft & Anomaly Detection

AT&C losses are tracked at utility level but identifying which specific meters cause those losses is manual, slow, and heavily dependent on individual inspector experience. A single theft case can go undetected for months. Tampered meters produce erratic readings that look like sensor noise. Peer-deviation cases — where a meter consistently draws 3–5x more than zone neighbors — are completely invisible without automated cross-meter comparison.

The gap is not data — BESCOM already has it. The gap is intelligence that converts raw 15-minute readings into actionable decisions.

3. System Architecture




Figure 1: End-to-end GridSense AI pipeline

Key design principle: Every flag that reaches an inspector comes with (1) a confidence score 0–1, (2) a plain-English explanation, (3) a SHAP feature breakdown, and (4) a full timestamped audit entry.

4. Implementation Details

4.1 Synthetic Data Generation (data_generator.py)

432,050 rows of realistic smart meter data across 50 meters, 5 zones, and 90 days (January–March 2024, Bangalore). All seeds fixed at 42 for full reproducibility.

- **Meter types:** 70% residential (sharp morning/evening peaks), 20% commercial (flat business-hours profile), 10% industrial (near-constant 24/7)
- **Realism:** Per-hour consumption multipliers, weekend boosts, Bangalore seasonal patterns (peak AC load April–May, coolest December), 8% Gaussian noise
- **Injected anomalies (20% of fleet — 10 meters):** THEFT (10–30% of normal after bypass date), TAMPER (8% chance per reading of 5–15x spike/drop), SUDDEN_DROP (1–5 day outage then recovery), PEER_DEVIATION (sustained 3–5x elevation vs neighbors)
- **Ground-truth labels:** is_anomaly column enables rigorous precision/recall evaluation

4.2 Feature Engineering (feature_engineering.py) — 25+ Features

Feature Group	Key Features	Primary Signal For
Lag Features	lag_1h, lag_24h, lag_48h, lag_7d, ratio_vs_24h, ratio_vs_7d	Theft detection
Rolling Stats	roll_mean/std 2h/24h/7d, zscore_24h, zscore_7d, cv_24h	Tamper detection
Peer Deviation	peer_ratio, peer_zscore, peer_diff (leave-one-out)	Theft & peer deviation
Rate of Change	delta_1step, delta_1h, accel_1step, abs_delta_1h	Sudden drop & theft onset
Cyclical Time	hour_sin/cos, dow_sin/cos, month_sin/cos, is_weekend	All patterns
Gap Features	gap_minutes, has_gap, cumulative_gaps	Data quality & tampering

Critical design: ratio_vs_24h = current / yesterday_same_time is the single strongest theft signal. Peer features use a **leave-one-out** zone mean so a thief cannot hide by inflating their own zone average.

4.3 Part A — Demand Forecasting (demand_forecast.py)

Model 1: Gradient Boosting Regressor — one model per zone (200 trees, max_depth=5, learning_rate=0.05, subsample=0.8). Zone-specific models capture each feeder's unique load profile. XGBoost-compatible: one-line swap.

Model 2: Seasonal Decomposition (Prophet-equivalent) — Fourier series representing daily (24h) and weekly (168h) demand cycles (fourier_order=5). Ridge regression finds optimal Fourier coefficients. Uncommenting 10 lines switches to Meta's Prophet.

Ensemble: Grid search over 0–100% blend ratios on held-out validation. Final = 0.70 x GradientBoosting + 0.30 x Seasonal. No test set contamination.

Zone Risk: load_factor = predicted / 95th_percentile_peak → LOW (<60%), MEDIUM (60–75%), HIGH (75–90%), CRITICAL (>90%). Maps directly to BESCOM's operational response levels.

Table 1: Demand Forecasting Results — Test Set (Days 76–90)

Model	MAE (kWh)	MAPE	R2
Naive Baseline (yesterday's value)	0.636	9.20%	0.887

Seasonal Model (Fourier)	0.597	9.56%	0.922
Gradient Boosting	0.004	0.07%	1.000
Ensemble (Final)	0.004	0.07%	1.000

99.3% improvement over naive baseline. On real BESCOM data expect production-grade MAPE of 3–8%.

4.4 Part B — Anomaly & Theft Detection ([anomaly_detector.py](#))

Model 1: Isolation Forest (200 trees, contamination=0.15) — anomalies are isolated in fewer random splits because they sit far from the dense cluster of normal data. Trained only on clean meters: learns what normal looks like, flags deviation.

Model 2: Autoencoder via MLPRegressor (20→64→16→64→20). The 16-neuron bottleneck forces compression of normal patterns. At inference, anomalous patterns produce high reconstruction error — threshold = 95th percentile of training reconstruction error.

Ensemble: 0.55 x IsolationForest + 0.45 x Autoencoder. IF catches global outliers; AE catches contextual anomalies ('this value is normal at 7 PM but anomalous at 3 AM').

4.4.1 Anomaly Type Classifier — Rule-Based, Fully Transparent

Type	Detection Rule
THEFT	ratio_vs_24h < 0.35 AND peer_ratio < 0.40 AND cv_24h < 0.5 (consistently low, not erratic)
TAMPER	cv_24h > 0.8 OR (abs_delta_1h > 2 x roll_mean_24h AND cv_24h > 0.4) (erratic, wild swings)
SUDDEN_DROP	consumption < 0.05 AND abs_delta_1h > 0.2 AND peer_ratio < 0.15 (sharp single event)
PEER_DEV	peer_ratio > 2.5 (consistently 2.5x above zone neighbors)

Rules are intentionally transparent — BESCOM engineers can read, audit, and update without retraining any ML model.

Table 2: Anomaly Detection Results — Test Set (Days 76–90)

Model	Precision	Recall	F1	ROC-AUC
Isolation Forest	0.27	0.48	0.35	0.82
Autoencoder	0.52	0.90	0.66	0.94
Ensemble (Raw)	0.11	1.00	0.19	0.95
Ensemble + FP Filter	0.16	1.00	0.28	0.95

FP filter suppresses 83.7% of raw flags before reaching inspectors. Ensemble achieves Recall = 1.00 — no theft case is missed.

4.5 Explainability & Audit Layer ([explainability.py](#))

- **SHAP Explainer:** Shapley values per flag — output: 'Flag driven by: low ratio_vs_24h (-0.89), low peer_ratio (-0.82)'. Top 5 features shown per alert.
- **Confidence Scorer:** Converts raw IsolationForest scores to 0–1 with labels: LOW / MEDIUM / HIGH.
- **False Positive Filter:** 4 named rules — (1) confidence < 0.55, (2) only 1 consecutive flag, (3) known maintenance window, (4) weak SHAP signal. Every suppression logged — nothing silently dropped.
- **Audit Logger:** Appends every decision to JSONL file — timestamp, meter_id, SHAP values, confidence, filter decision, reason. Append-only, grep-friendly, full regulatory compliance trail.

4.6 Streamlit Dashboard (dashboard.py)

- **Tab 1 — Demand & Risk:** Glassmorphic zone risk cards (pulse animation for CRITICAL zones), toggle between AI Insight and Pie Chart views, actual vs predicted demand chart with CRITICAL period shading.
- **Tab 2 — Anomaly Alerts:** Real-time inspector alert feed sorted by confidence score, anomaly type breakdown chart, inspector action queue (top 5 meters).
- **Tab 3 — Explainability:** Global feature importance chart with plain-English captions for non-technical BESCOM staff.

Fully Dockerized and deployed on HuggingFace Spaces. All loaders use @st.cache_data for fast in-memory interaction.

5. Non-Negotiables Compliance

Requirement	How GridSense AI Satisfies It
No modification to existing systems	Pure read-only overlay on MDMS CSV exports — no API changes needed
Works as decision-support layer	Every flag requires human inspector confirmation — zero automated actions
Uses masked/synthetic data only	All dev and demo uses fully synthetic data — no real BESCOM data anywhere
Outputs explainable and auditable	SHAP per flag + plain-English summaries + full JSONL audit trail
False positives minimised and visible	83.7% suppression rate; every suppression has a named reason on record
No hosted LLM on sensitive data	Zero external API calls — all inference local, no data leaves the machine

6. Project File Structure

File	Purpose	Input → Output
data_generator.py	Synthetic data generation	— → smart_meter_data.csv
feature_engineering.py	25+ feature pipeline	smart_meter_data.csv → features_meter.csv, features_zone.csv
demand_forecast.py	Part A: demand forecasting	features_zone.csv → forecasts.csv, zone_risk.csv
anomaly_detector.py	Part B: anomaly detection	features_meter.csv → anomaly_results.csv
explainability.py	SHAP, FP filter, audit log	anomaly_model_if.pkl → audit_log.jsonl
dashboard.py	Streamlit interactive dashboard	All CSVs → live UI
requirements.txt	Python dependencies	—
Dockerfile	Container for HF deployment	—

7. Technology Stack

Category	Technologies
Core Language	Python 3.13
ML Models	scikit-learn: GradientBoostingRegressor, IsolationForest, MLPRegressor, Ridge Regression

Explainability	SHAP (TreeExplainer for IsolationForest)
Dashboard	Streamlit, Plotly (interactive charts)
Data Processing	pandas, numpy, joblib (model serialisation)
Deployment	Docker, HuggingFace Spaces
Future (Final)	XGBoost, LightGBM, PyTorch LSTM, Meta Prophet, Apache Kafka, Dask

8. Final Round Roadmap

Area	Planned Enhancement	Timeline
Real Data	Live MDMS API adaptor, sub-5-minute latency from reading to dashboard	Final Round
Models	PyTorch LSTM Autoencoder for sequential time-series, XGBoost/LightGBM	Final Round
Forecasting	Conformal Prediction intervals for calibrated uncertainty bounds, Meta Prophet	Final Round
Geography	GIS feeder heatmap on Bangalore topology, transformer-level aggregation	Final Round
Feedback	Inspector outcome capture (Confirmed/FP) feeds quarterly FP filter retraining	Post-Pilot
Scale	Dask + Apache Kafka for real-time streaming across millions of meters	Phase 2

- **PyTorch LSTM Autoencoder:** Replace MLPRegressor with an LSTM-based autoencoder for true sequential time-series learning — capturing temporal dependencies across days and weeks.
- **Conformal Prediction:** Calibrated uncertainty bounds on forecasts — '42 kWh $\pm$ 3.2 at 90% confidence' — rather than a single point estimate.
- **GIS Feeder Heatmap:** Overlay zone risk on Bangalore's actual feeder topology. Spatial theft cluster visualization — multiple neighbouring meters dropping simultaneously is a stronger signal than any individual.
- **Inspector Feedback Loop:** Outcome capture (Confirmed Theft / False Positive) feeds quarterly retraining of FP filter rules. Persistent meter-level risk scores accumulate across time.

9. Impact Potential & Closing Statement

A conservative estimate based on published AT&C loss data: if GridSense AI helps BESCOM identify and act on even 30% of active theft cases among its smart meter fleet, the revenue recovery potential runs into hundreds of crores per year. Beyond revenue, proactive demand forecasting reduces emergency load shedding events, extends transformer lifespan by avoiding sustained overload, and gives grid operations the reaction time they currently lack.

GridSense AI is not a research prototype. It is a deployable, explainable, locally-running decision layer that sits on top of infrastructure BESCOM already has, requires no new hardware, produces no external data exposure, and delivers value from day one of pilot deployment. The full pipeline runs end-to-end in under 2 minutes on a single machine with 8 GB RAM.

Live Demo	https://huggingface.co/spaces/rishav123987/gridsense-ai-smart-meter-dashboard
Source Code	https://github.com/lokesh07kewat/BESCOM-Smart-Intelligence

All data is synthetic. No real BESCOM customer data is used or stored. All inference runs locally.